

# Temporal State Models

Represent the ‘world’ as a set of random variables  $\mathbf{X}$

$\mathbf{X} = \{x, y\}$       location on the ground plane

$\mathbf{X} = \{x, y, z\}$       position in the 3D world

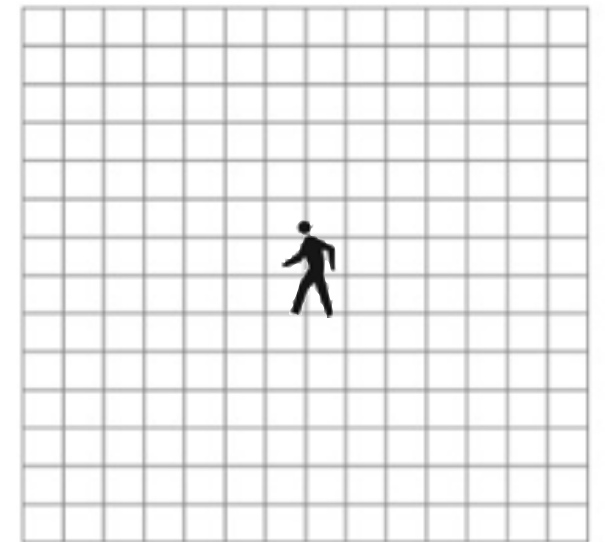
$\mathbf{X} = \{x, \dot{x}\}$       position and velocity

$\mathbf{X} = \{x, \dot{x}, f_1, \dots, f_n\}$   
position, velocity and  
location of landmarks

## Object tracking (localization)

$$\mathbf{X} = \{x, y\}$$

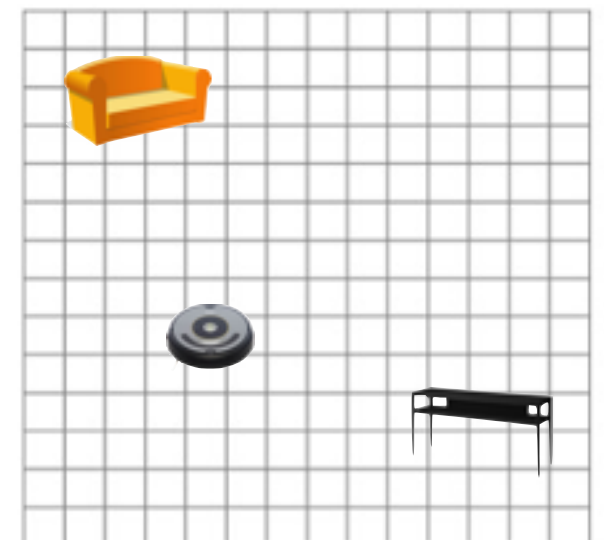
e.g., location on the ground plane



## Object location and world landmarks (localization and mapping)

$$\mathbf{X} = \{x, \dot{x}, f_1, \dots, f_n\}$$

e.g., position and velocity of  
robot and location of landmarks



$\mathbf{X}_t$




The state of the world changes over time

$$\mathbf{X}_t$$

The state of the world changes over time

So we use a sequence of random variables:

$$\mathbf{X}_0, \mathbf{X}_1, \dots, \mathbf{X}_t$$

$$\mathbf{X}_t$$


The state of the world changes over time

So we use a sequence of random variables:

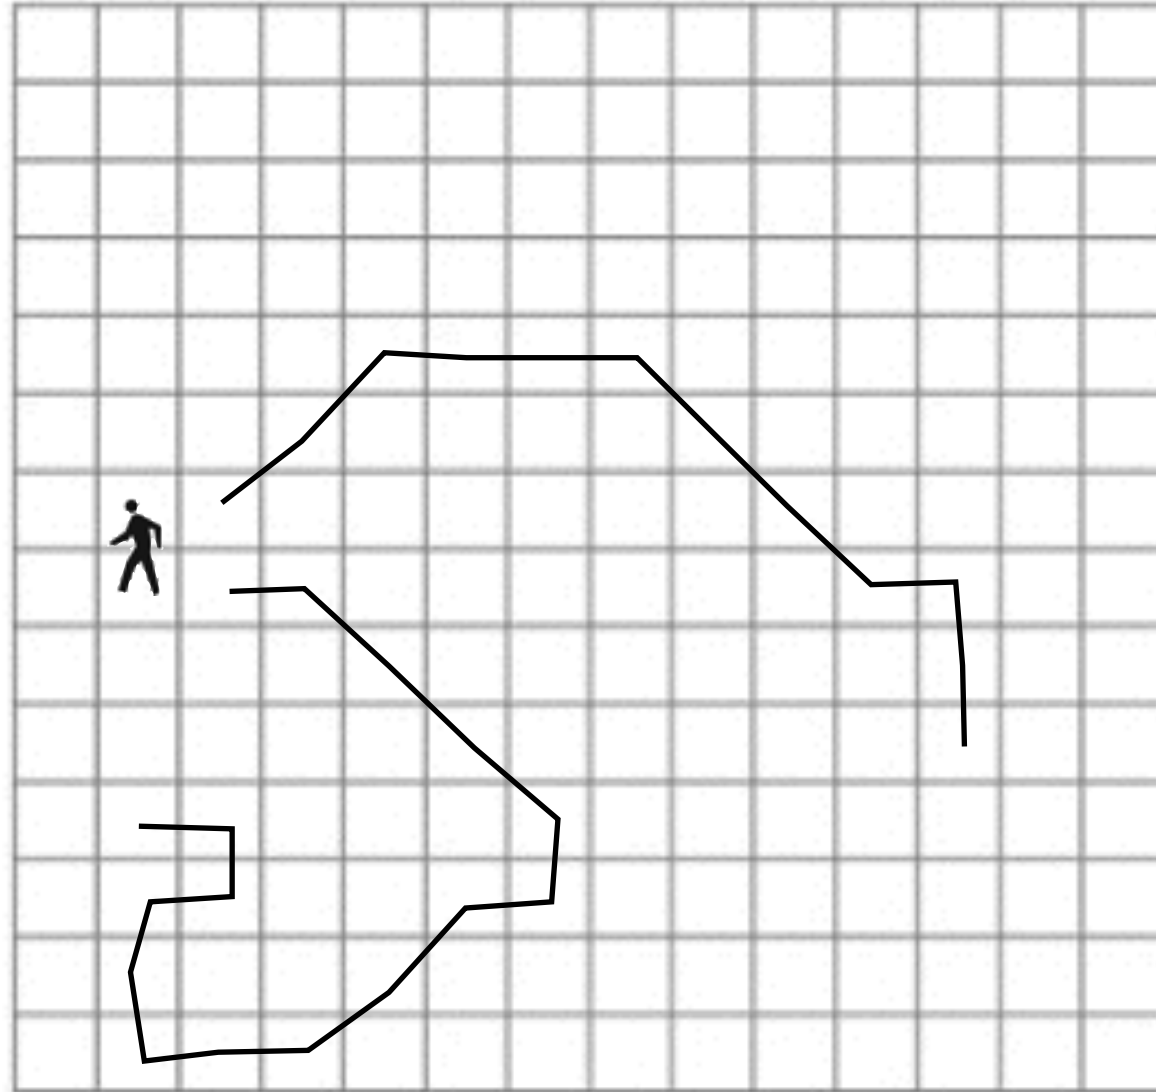
$$\mathbf{X}_0, \mathbf{X}_1, \dots, \mathbf{X}_t$$

The state of the world is usually **uncertain** so we think in terms of a distribution

$$P(\mathbf{X}_0, \mathbf{X}_1, \dots, \mathbf{X}_t)$$

*How big is the space of this distribution?*

If the state space is  $\mathbf{X} = \{\mathbf{x}, \mathbf{y}\}$  the location on the ground plane



$$P(\mathbf{X}_0, \mathbf{X}_1, \dots, \mathbf{X}_t)$$

is the probability over all possible trajectories through a room of length  $t+1$

When we use a sensor (camera),  
we don't have direct access to the state but noisy  
observations of the state

$$\mathbf{E}_t$$

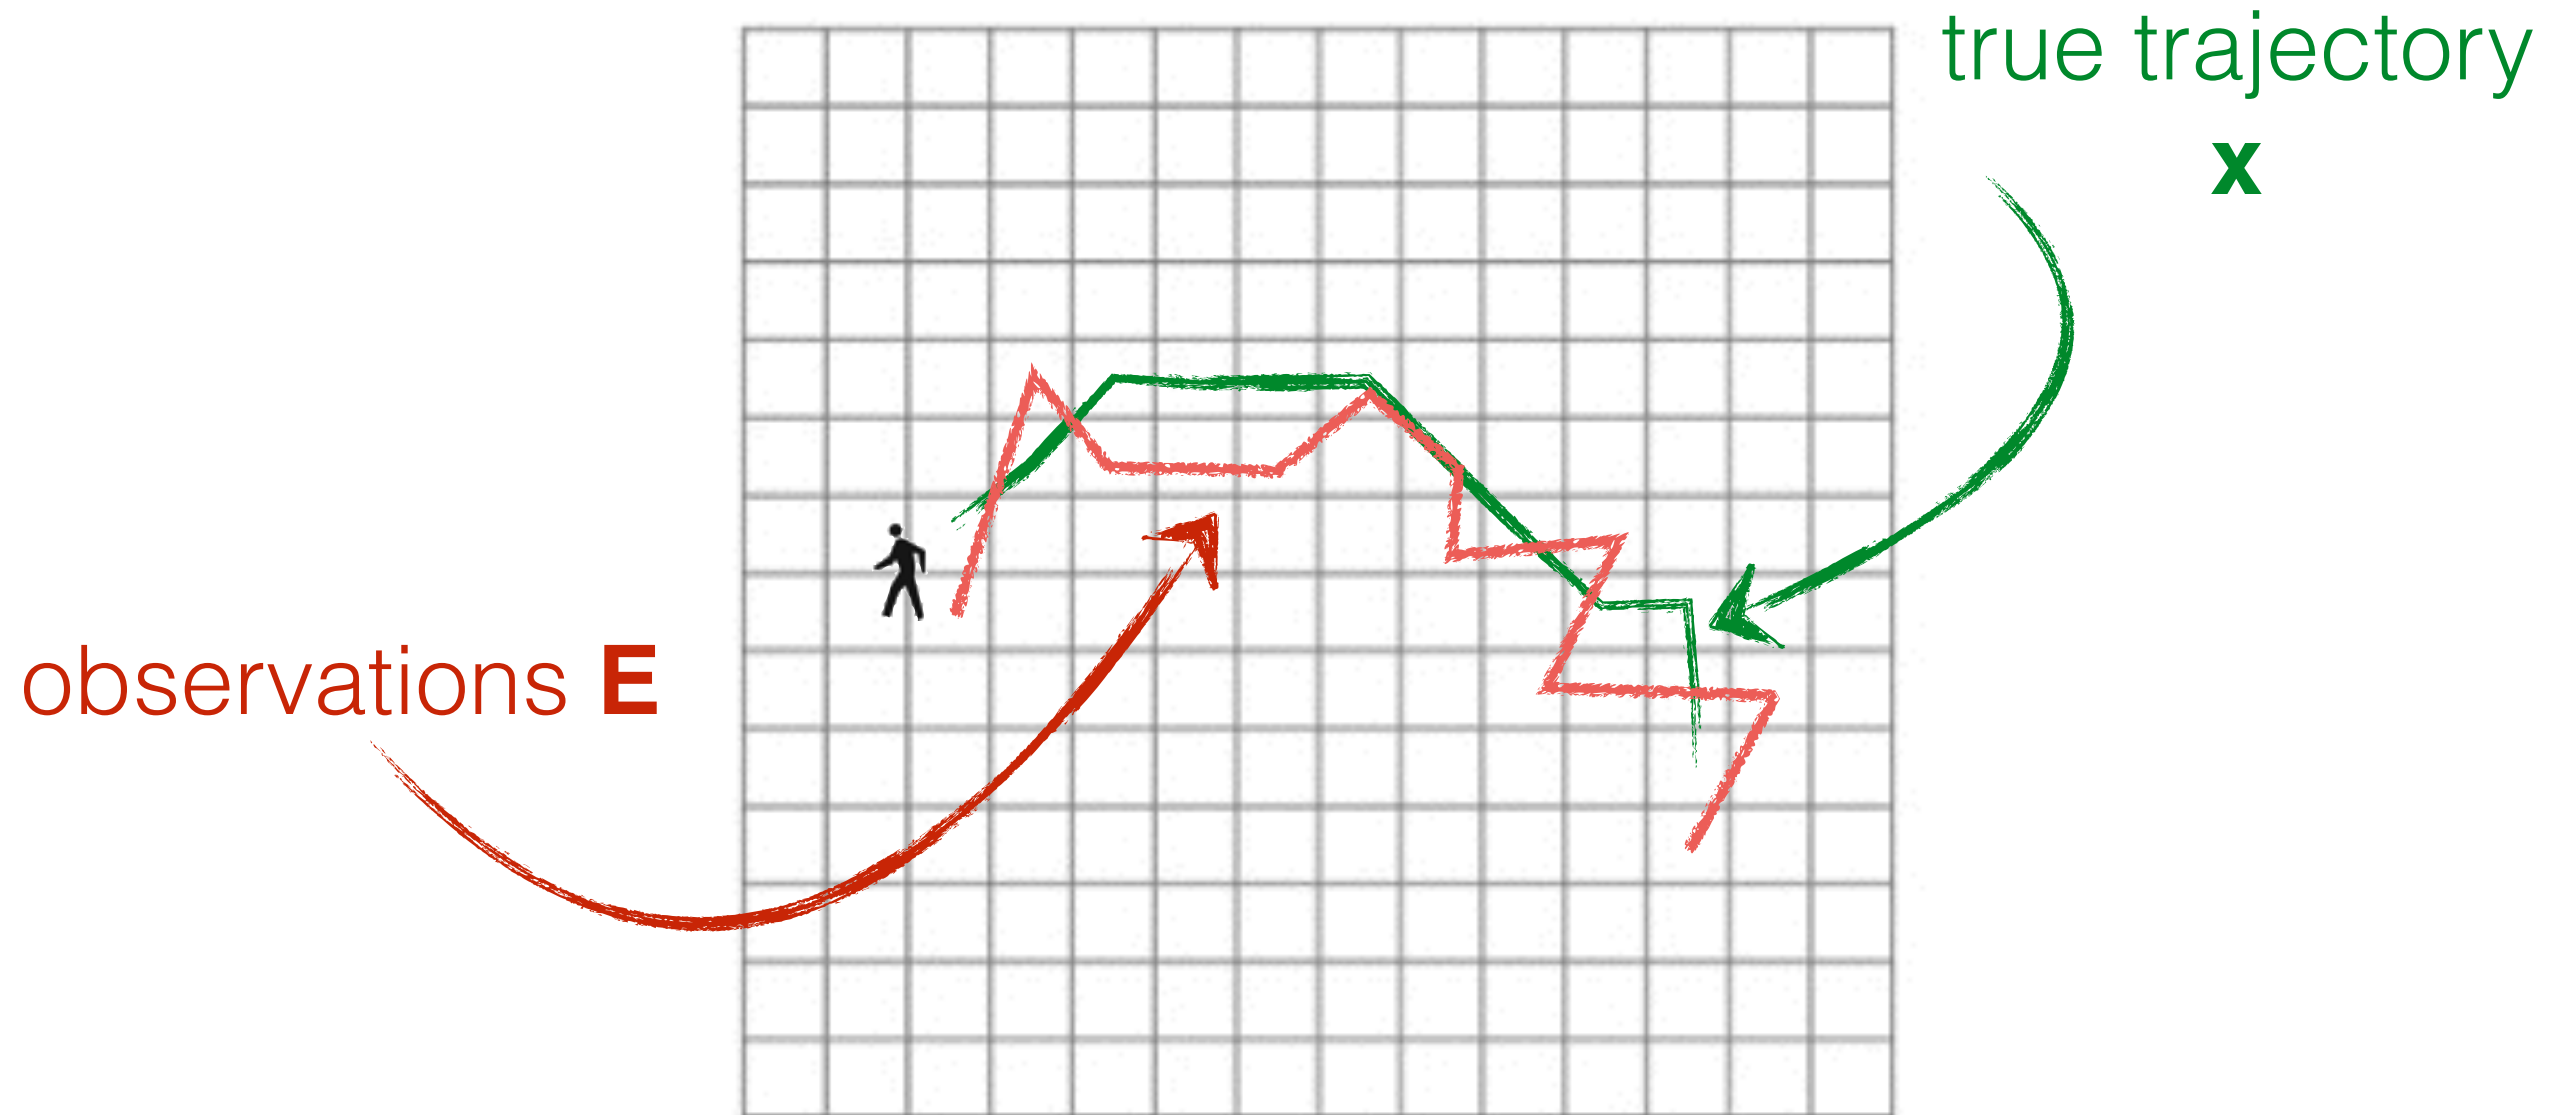
$$\mathbf{X}_0, \mathbf{X}_1, \dots, \mathbf{X}_t, \mathbf{E}_1, \mathbf{E}_2, \dots, \mathbf{E}_t$$

(**all** possible ways of observing **all** possible trajectories)

*How big is the space of this distribution?*



**all** possible ways of observing **all** possible trajectories of length  $t$



So we think of the world in terms of the distribution

$$P(\underbrace{X_0, X_1, \dots, X_t}_{\substack{\text{unobserved variables} \\ \text{(hidden state)}}, \underbrace{E_1, E_2, \dots, E_t}_{\substack{\text{observed variables} \\ \text{(evidence)}}})$$

So we think of the world in terms of the distribution

$$P(\underbrace{X_0, X_1, \dots, X_t}_{\substack{\text{unobserved variables} \\ \text{(hidden state)}}, \underbrace{E_1, E_2, \dots, E_t}_{\substack{\text{observed variables} \\ \text{(evidence)}}})$$

*How big is the space of this distribution?*

So we think of the world in terms of the distribution

$$P(\underbrace{X_0, X_1, \dots, X_t}_{\substack{\text{unobserved variables} \\ \text{(hidden state)}}, \underbrace{E_1, E_2, \dots, E_t}_{\substack{\text{observed variables} \\ \text{(evidence)}}})$$

*How big is the space of this distribution?*

*Can you think of a way to reduce the space?*

## Reduction 1. Stationary process assumption:

*‘a process of change that is governed by laws that do not themselves change over time.’*

$$P(\mathbf{E}_t | \mathbf{X}_t) = P_t(\mathbf{E}_t | \mathbf{X}_t)$$



the model doesn't change over time

## Reduction 1. Stationary process assumption:

*‘a process of change that is governed by laws that do not themselves change over time.’*

$$P(\mathbf{E}_t | \mathbf{X}_t) = P_t(\mathbf{E}_t | \mathbf{X}_t)$$



the model doesn't change over time

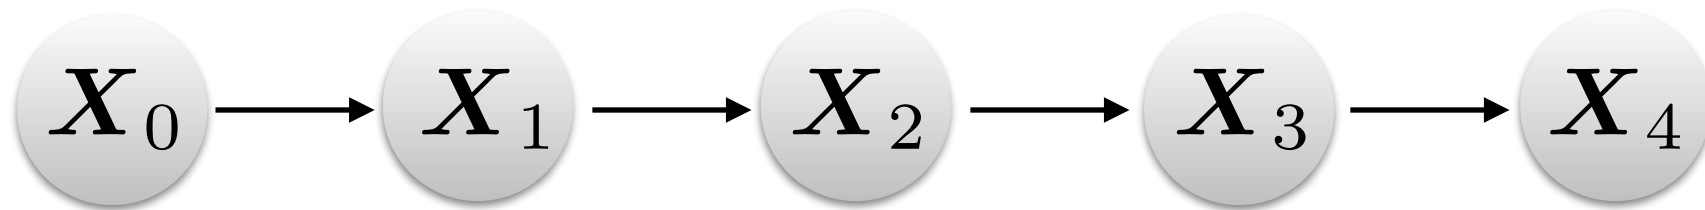
Only have to store **one** model.

*Is this a reasonable assumption?*

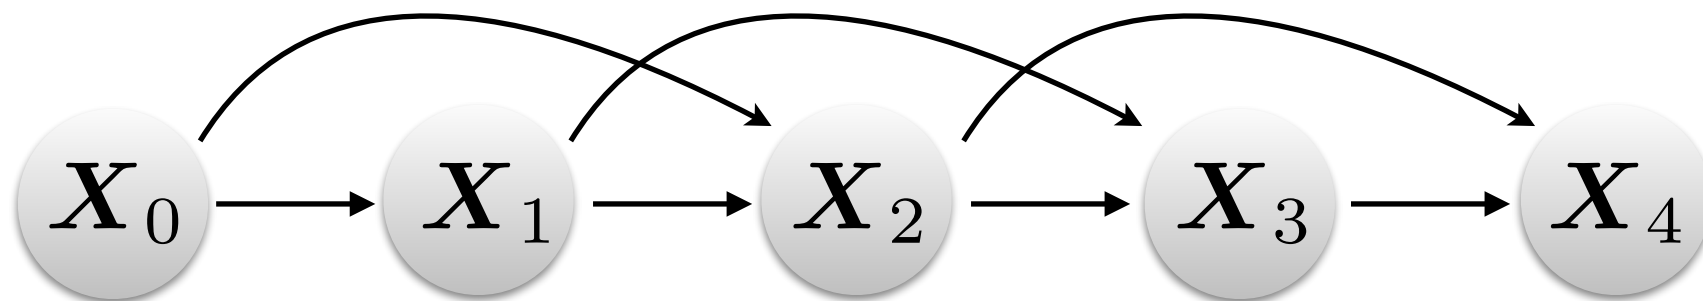
## Reduction 2. Markov Assumption:

*'the current state only depends on a finite history of previous states.'*

First-order Markov Model:  $P(\mathbf{X}_t | \mathbf{X}_{t-1})$ .



Second-order Markov Model:  $P(\mathbf{X}_t | \mathbf{X}_{t-1}, \mathbf{X}_{t-2})$



(this relationship is called the **motion** model)

## Reduction 2. Markov Assumption:

*‘the current observation only depends on current state.’*

The current observation is usually most influenced by the current state

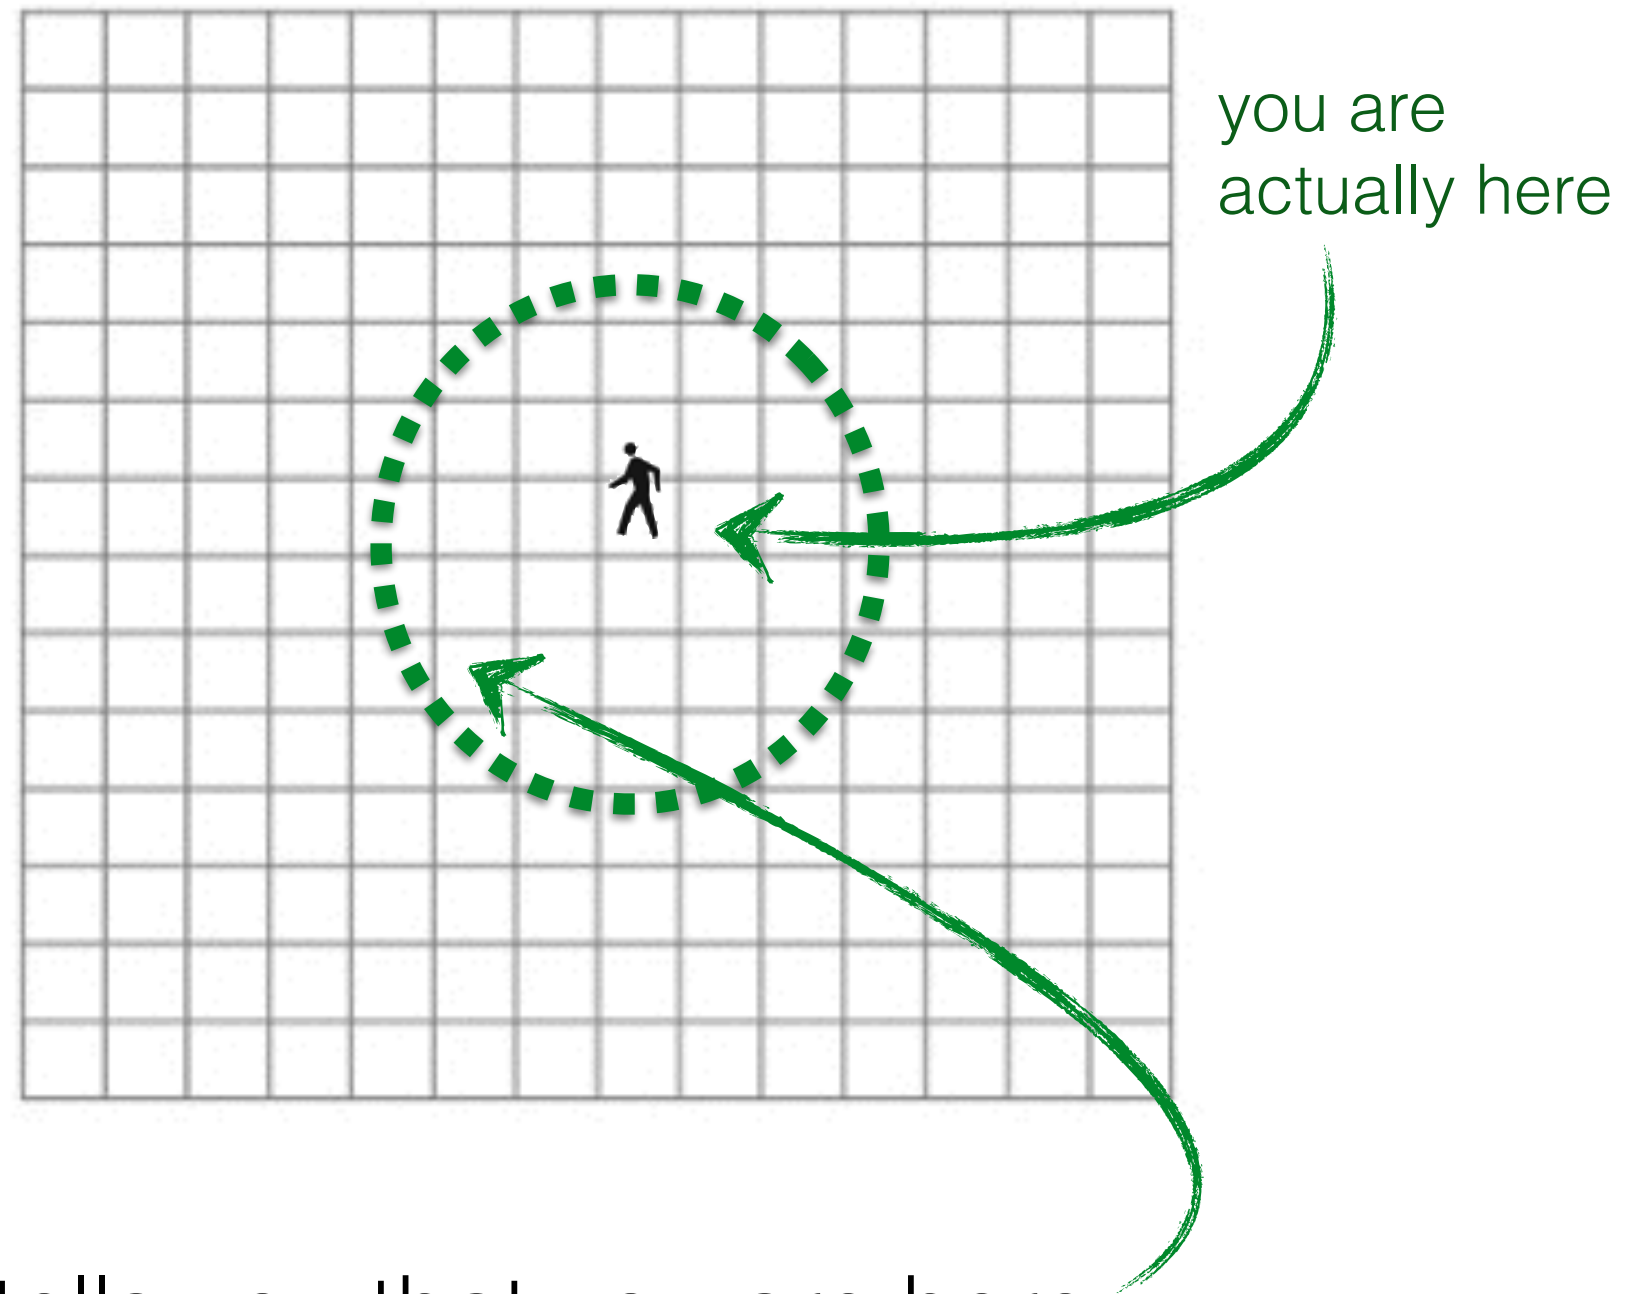
$$P(\mathbf{E}_t | \mathbf{X}_t)$$

(this relationship is called the **observation** model)

*Can you think of an observation of a state?*



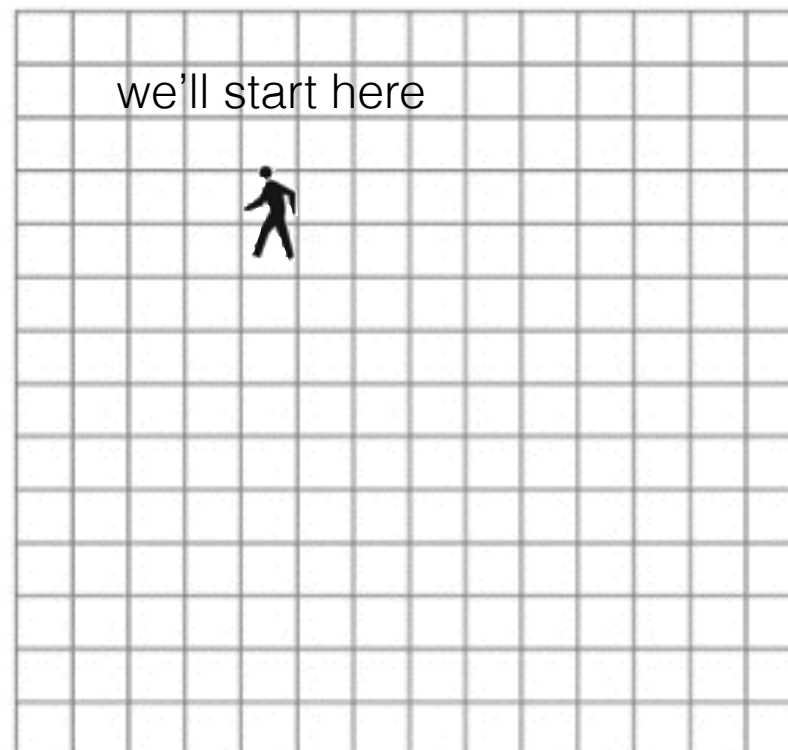
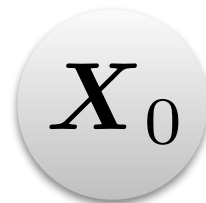
For example, GPS is a noisy observation of location.



But GPS tells you that you are here  
with probability  $P(\mathbf{E}_t | \mathbf{X}_t)$

## Reduction 3. Prior State Assumption:

*'we know where the process (probably) starts'*



## Joint Probability of a Temporal Sequence

$$P(\mathbf{X}_0) \prod_{t=1}^T P(\mathbf{X}_t | \mathbf{X}_{t-1}) P(\mathbf{E}_t | \mathbf{X}_t)$$

state prior  
prior

motion model  
transition model

sensor model  
observation model

## Joint Probability of a Temporal Sequence

$$P(\mathbf{X}_0) \prod_{t=1}^T P(\mathbf{X}_t | \mathbf{X}_{t-1}) P(\mathbf{E}_t | \mathbf{X}_t)$$

state prior  
prior

motion model  
transition model

sensor model  
observation model

## Joint Distribution for a Dynamic Bayesian Network

specific instances of a DBN  
covered in this class

Hidden Markov Model

(typically taught as discrete but not necessarily)

Kalman Filter

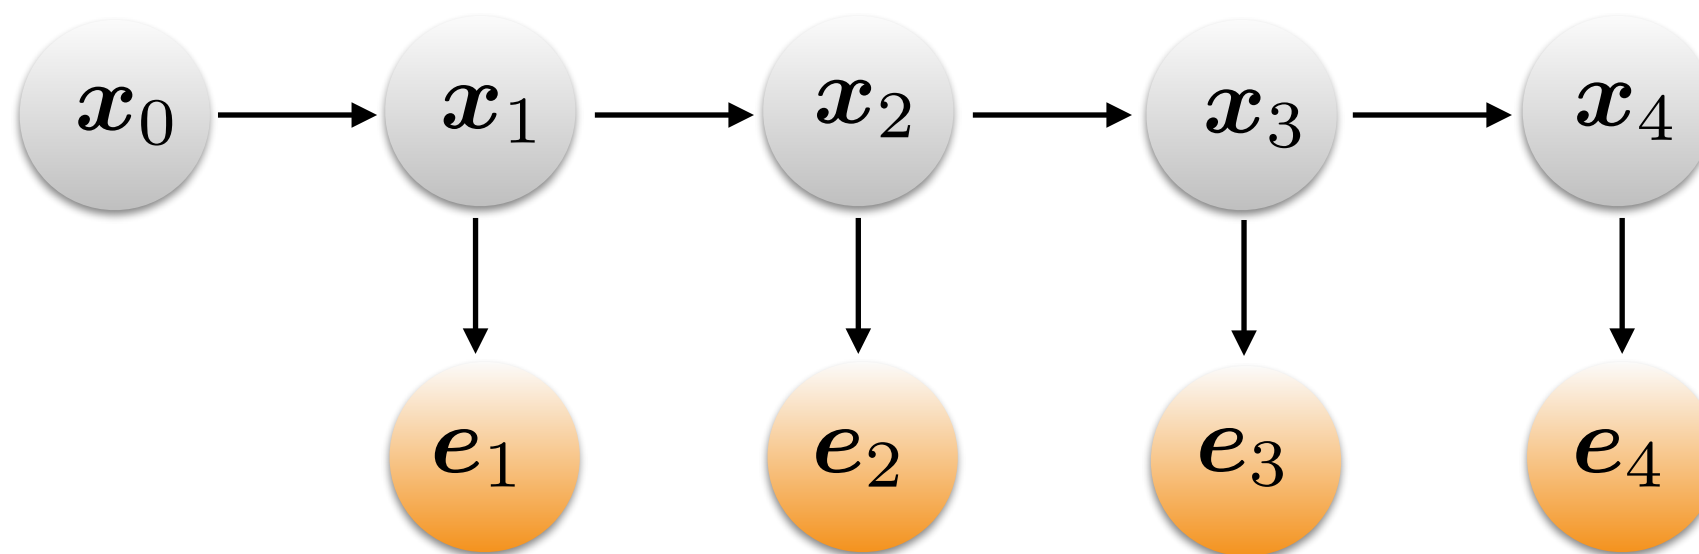
(Gaussian motion model, prior and observation model)



Kalman Filter

Examples up to now have been **discrete** (binary) random variables

Kalman 'filtering' can be seen as a special case of a temporal inference with continuous random variables



Everything is continuous...

$x$        $e$        $P(x_0)$        $P(e|x)$        $P(x_t|x_{t-1})$

probability distributions are no longer tables but functions

# Making the connection to the ‘filtering’ equations

# (Discrete) Filtering

## Tables

## Tables

## Tables

$$P(\mathbf{X}_{t+1}|\mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1}|\mathbf{X}_{t+1}) \sum_{\mathbf{X}_t} P(\mathbf{X}_{t+1}|\mathbf{X}_t) P(\mathbf{X}_t|\mathbf{e}_{1:t})$$

# Kalman Filtering

# Gaussian

# Gaussian

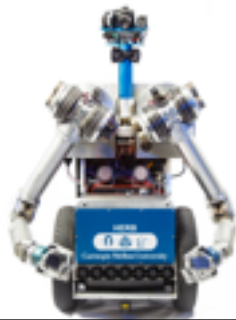
## Gaussian

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto \underbrace{P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1})}_{\text{observation model}} \int_{\mathbf{x}_t} \underbrace{P(\mathbf{X}_{t+1} | \mathbf{x}_t)}_{\text{motion model}} \underbrace{P(\mathbf{x}_t | \mathbf{e}_{1:t})}_{\text{belief}} d\mathbf{x}_t$$

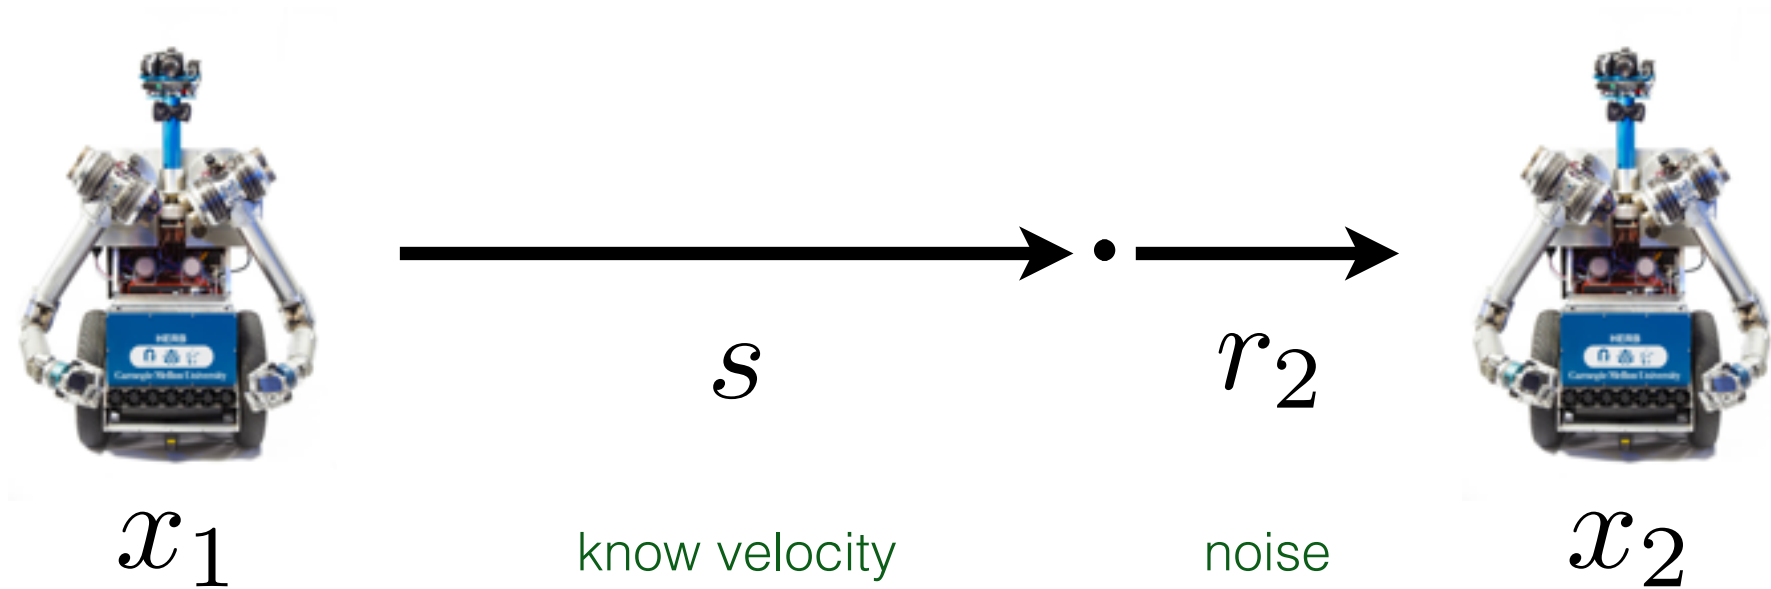
integral because  
continuous PDFs

Simple, 1D example...

$x$





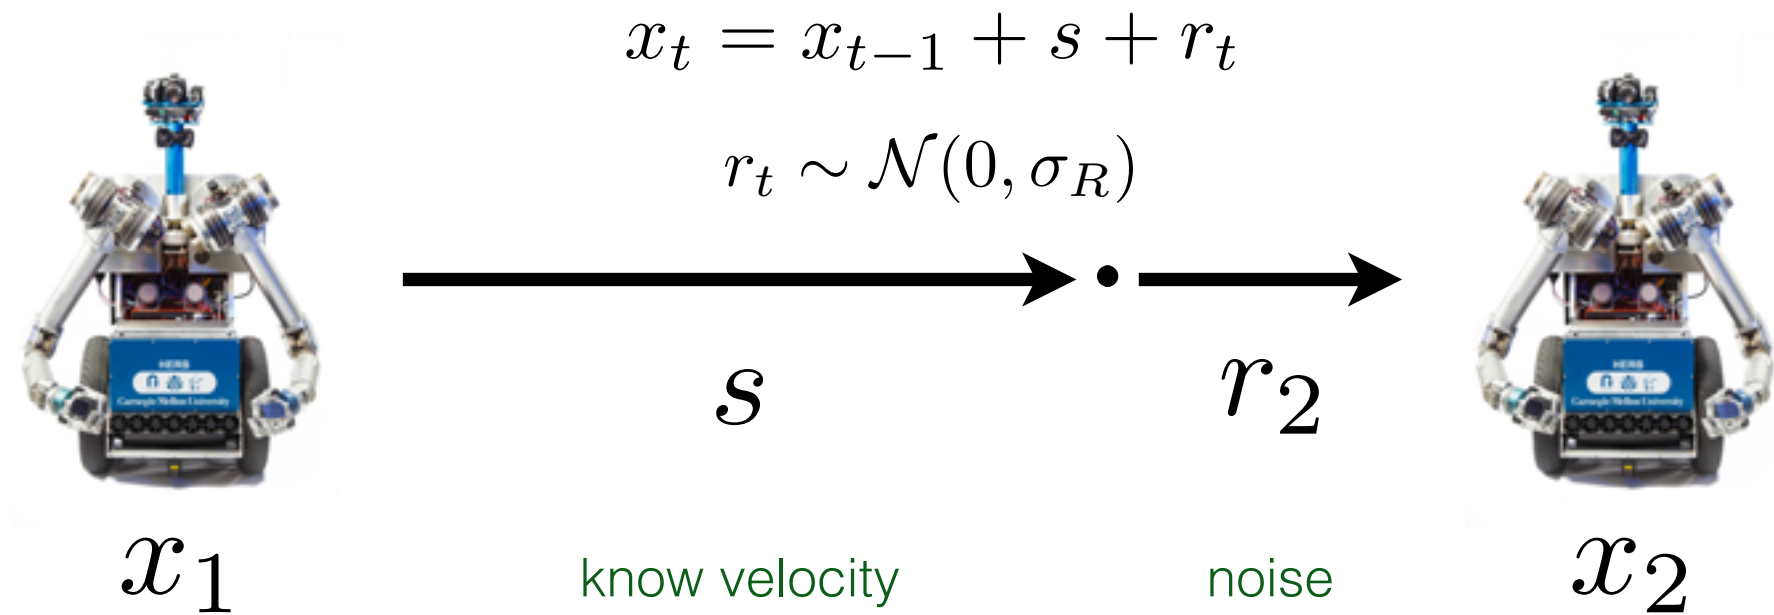


$$x_t = x_{t-1} + s + r_t$$

$$r_t \sim \mathcal{N}(0, \sigma_R)$$

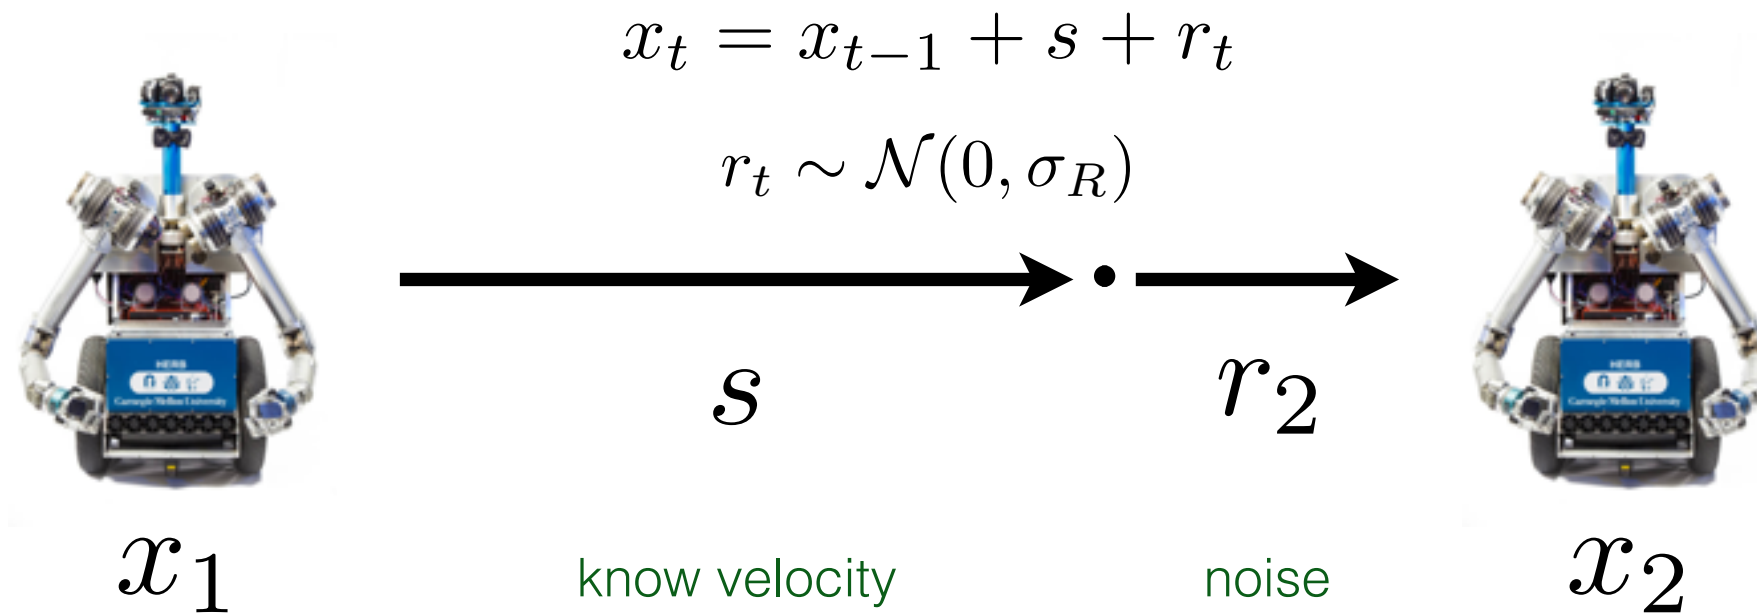
'sampled from'

# System (motion) model



*How do you represent the motion model?*

$$P(x_t | x_{t-1})$$



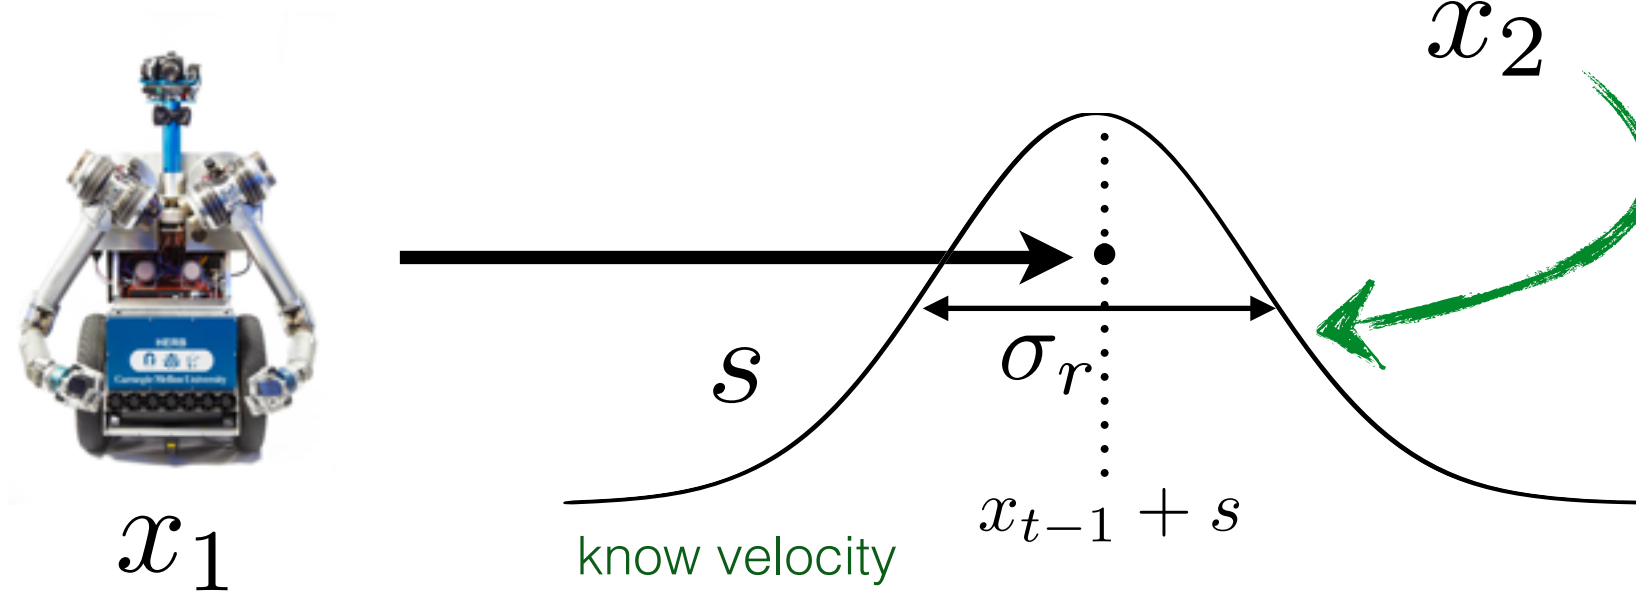
*How do you represent the motion model?*

A linear Gaussian (continuous) transition model

$$P(x_t | x_{t-1}) = \mathcal{N}(x_t; x_{t-1} + s, \sigma_r)$$

Labels:  $x_{t-1} + s$  (mean),  $\sigma_r$  (standard deviation)

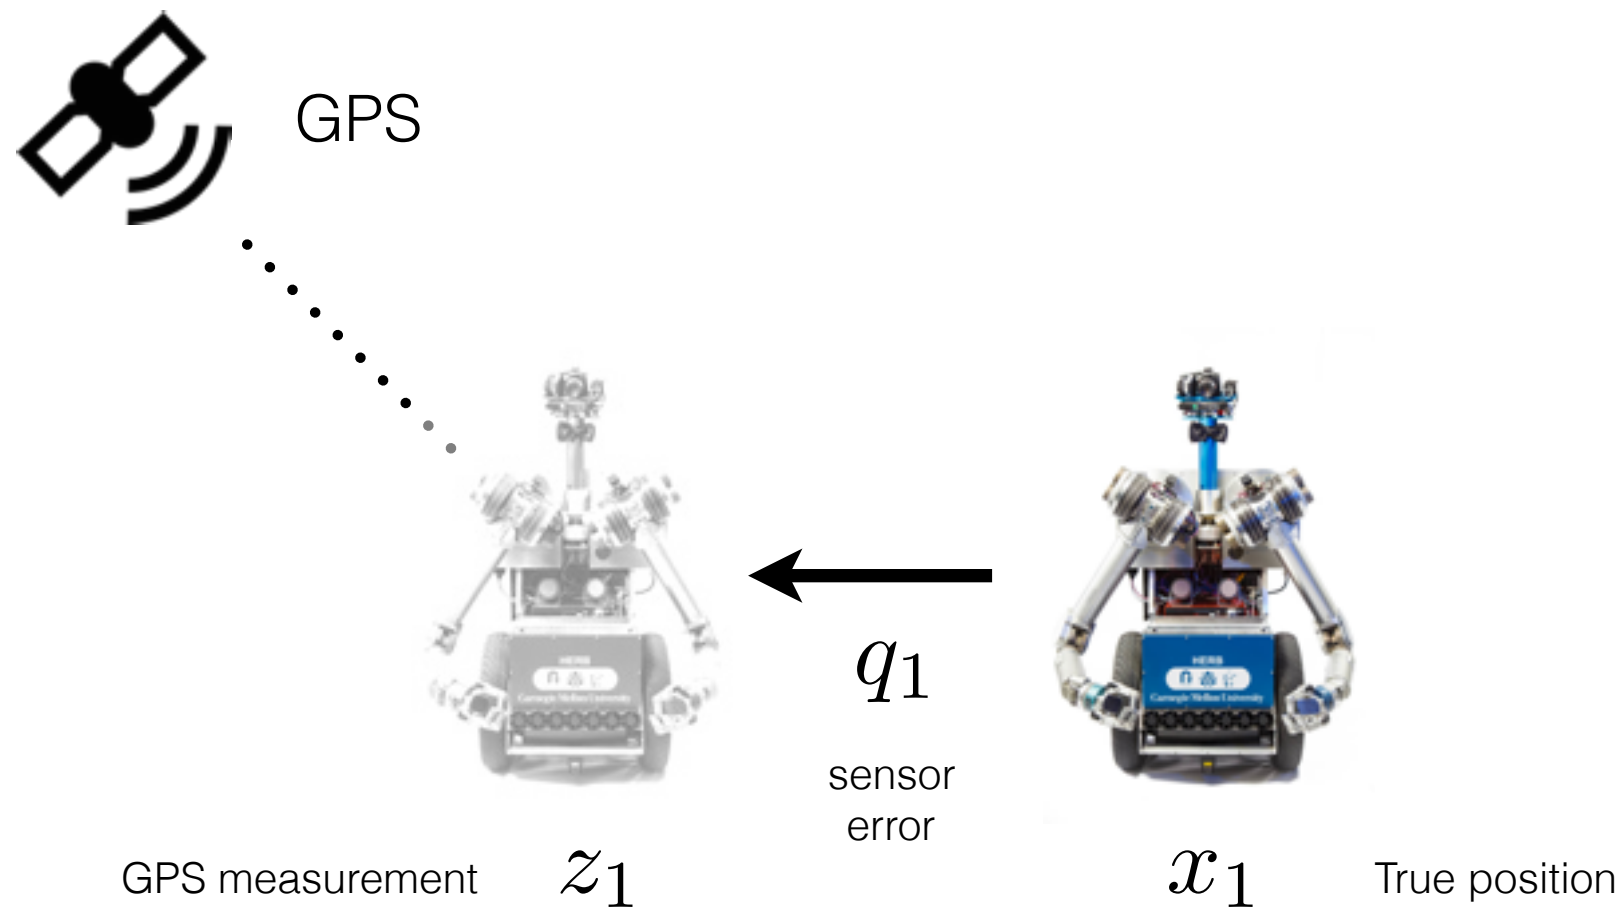
*How can you visualize this distribution?*



A linear Gaussian (continuous) transition model

$$P(x_t | x_{t-1}) = \mathcal{N}(x_t; x_{t-1} + s, \sigma_r)$$

*Why don't we just use a table as before?*

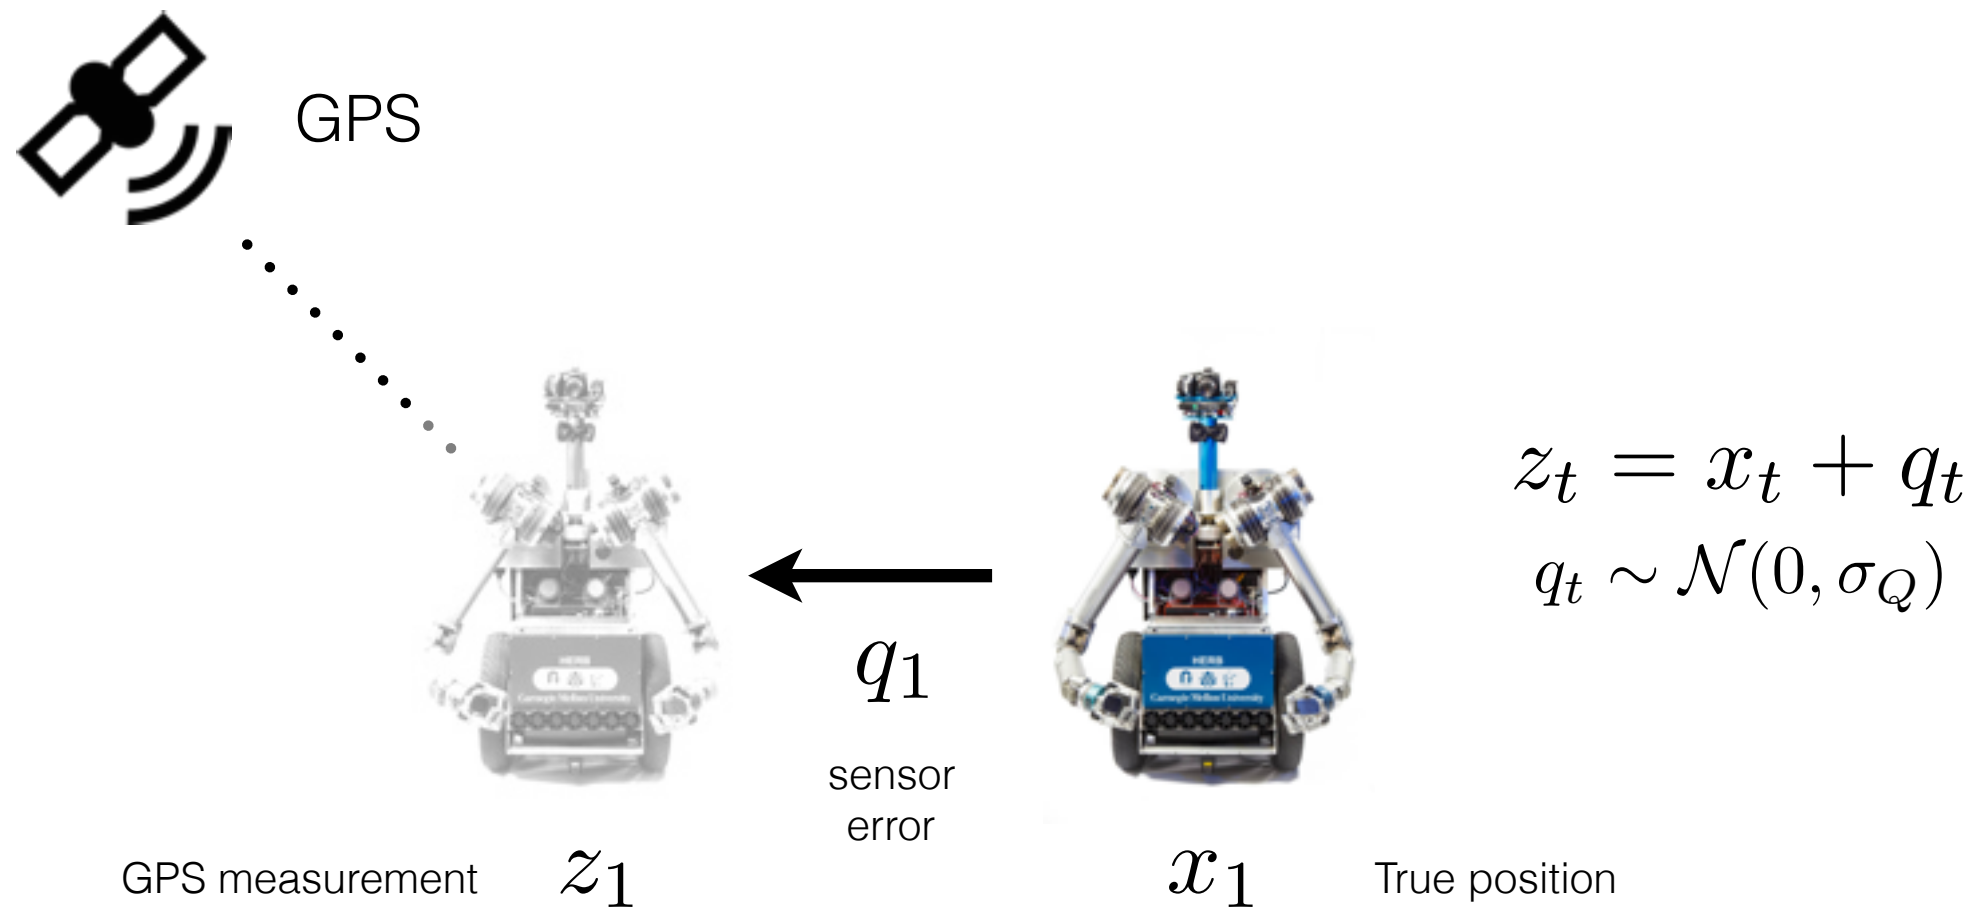


$$z_t = x_t + q_t$$

$$q_t \sim \mathcal{N}(0, \sigma_Q)$$

sampled from a Gaussian

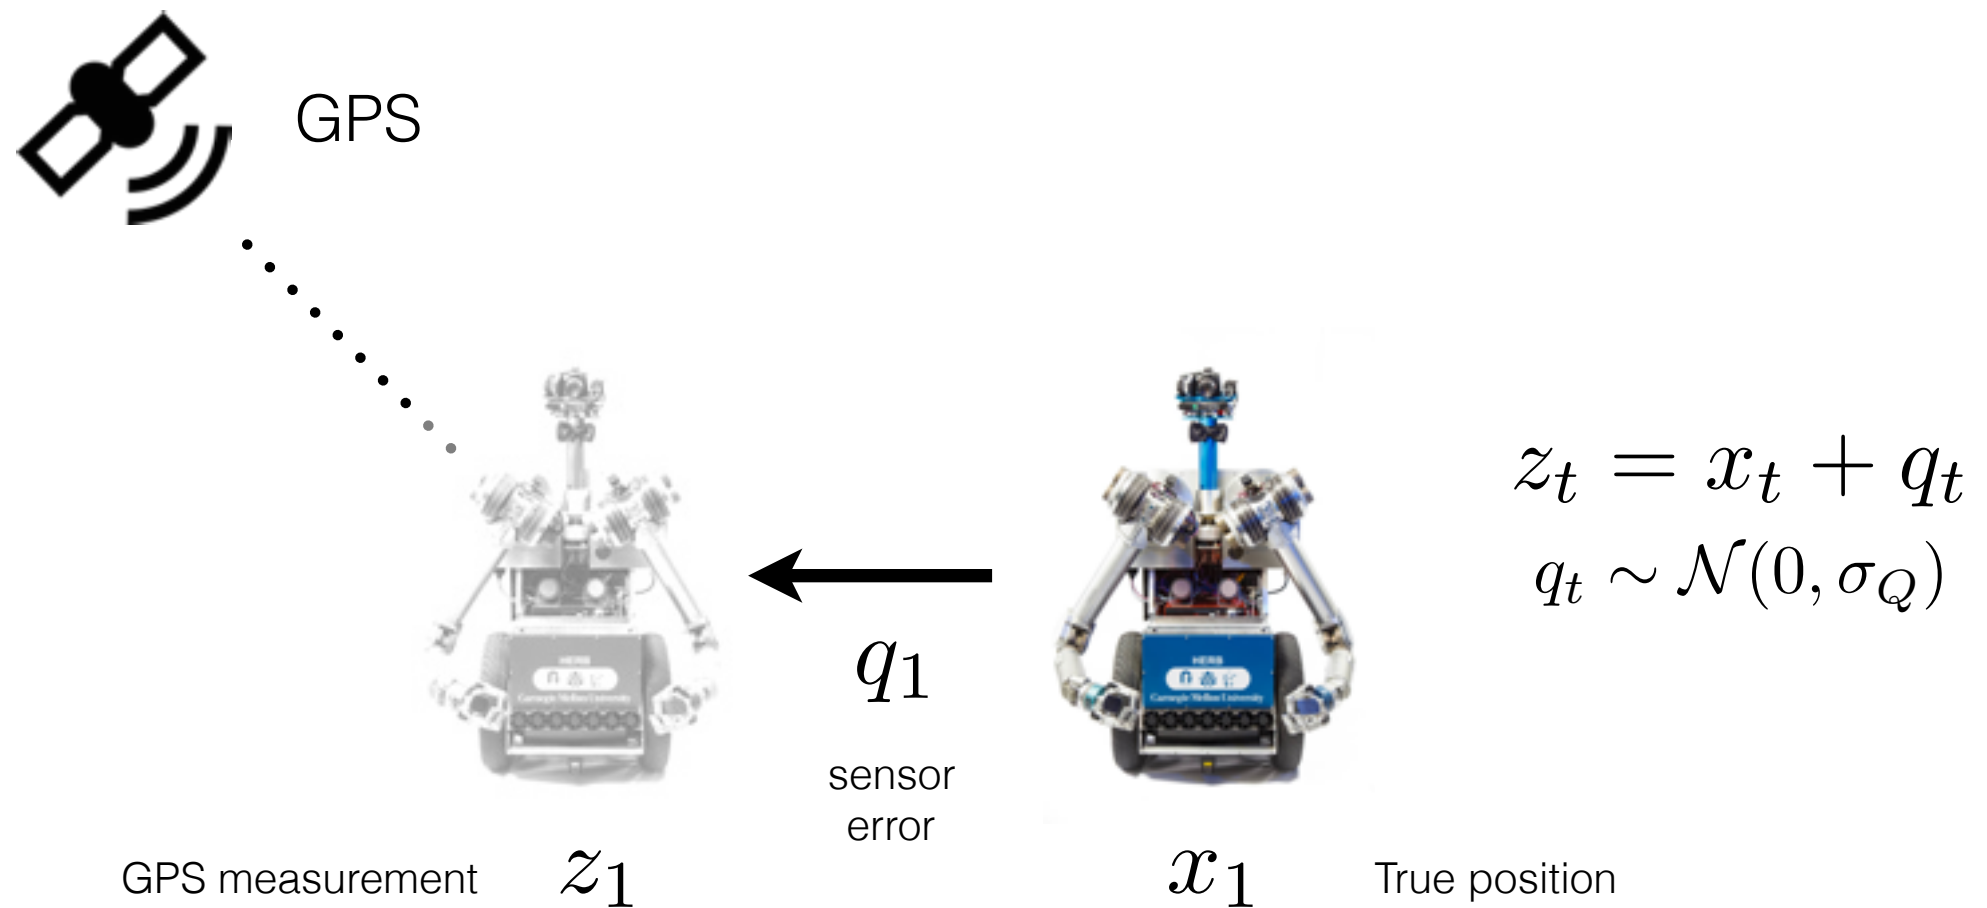
# Observation (measurement) model



*How do you represent the observation (measurement) model?*

$$P(e|x)$$

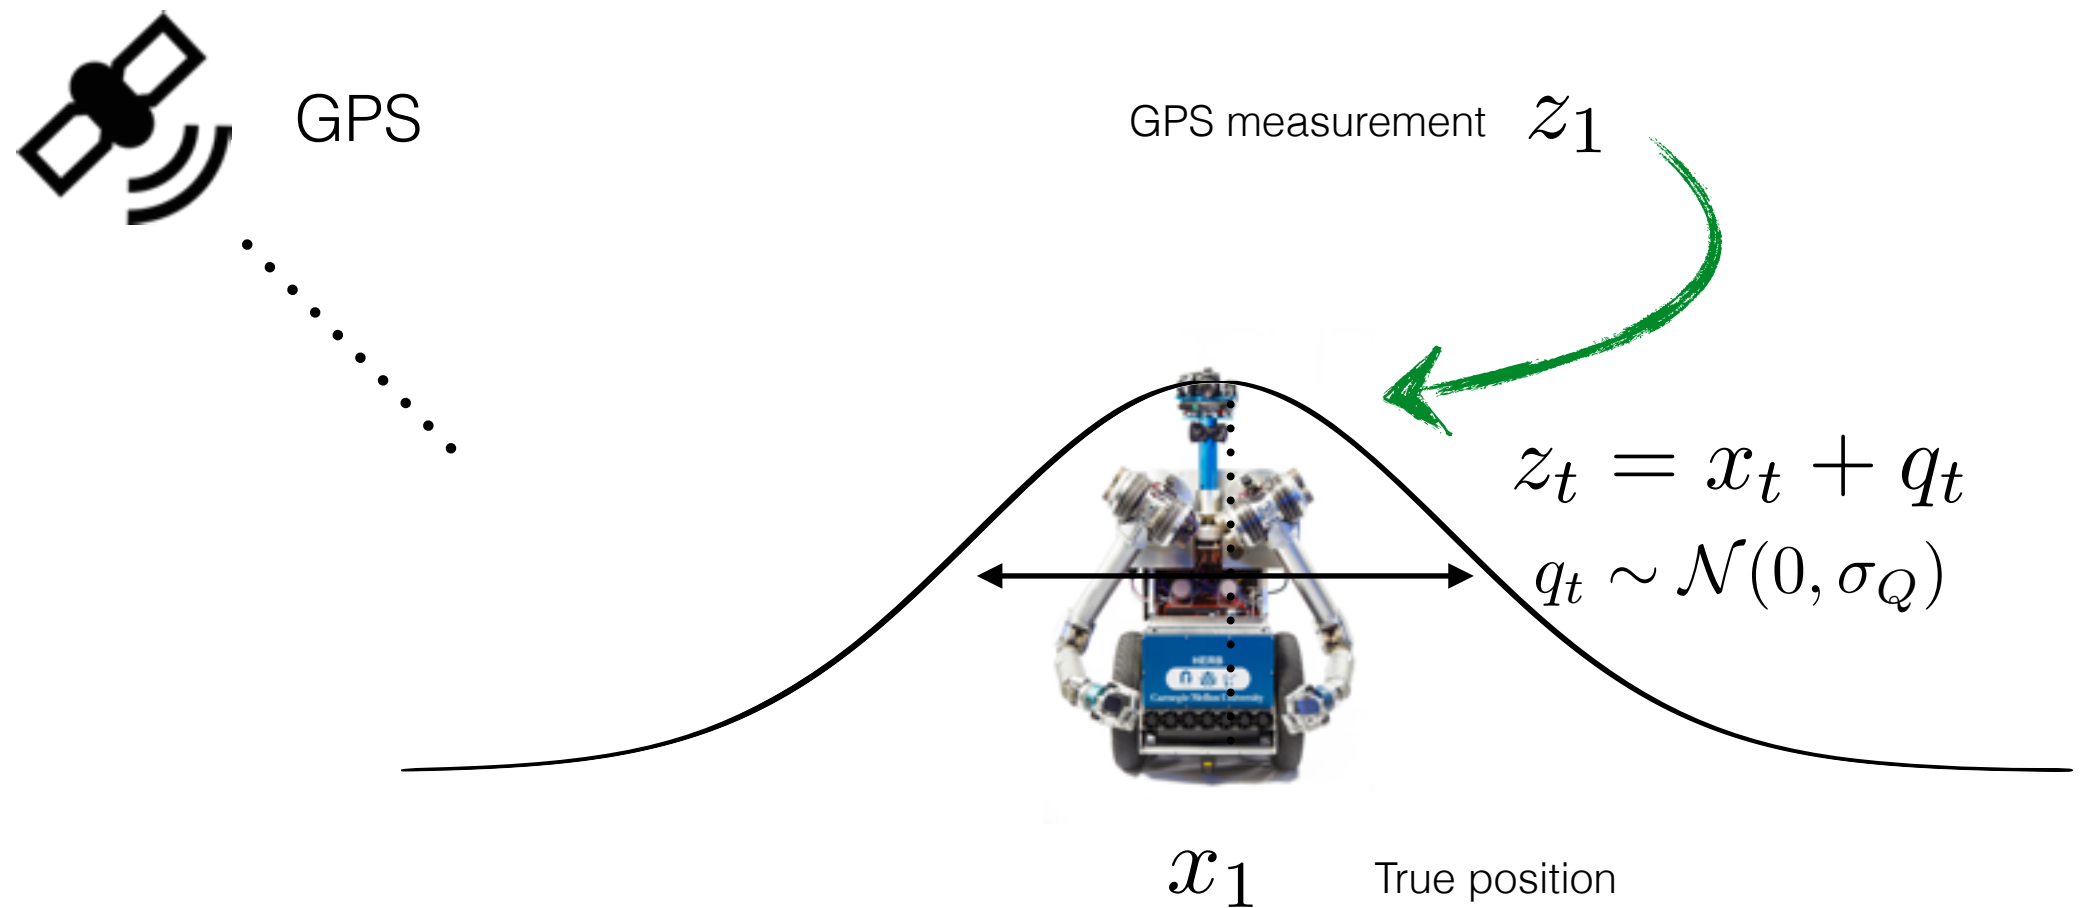
e represents z



*How do you represent the observation (measurement) model?*

Also a linear Gaussian model

$$P(z_t | x_t) = \mathcal{N}(z_t; x_t, \sigma_Q)$$

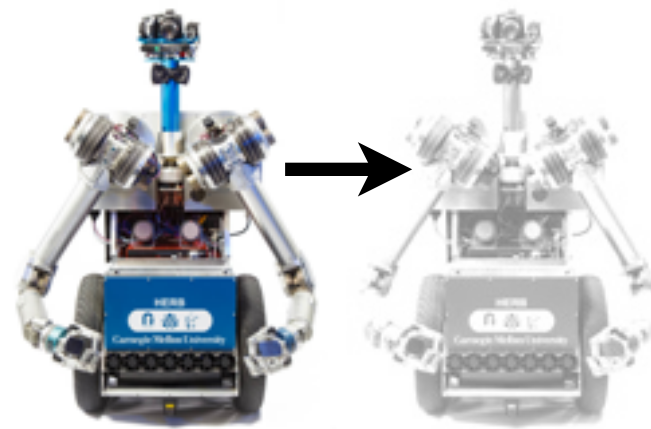


*How do you represent the observation (measurement) model?*

Also a linear Gaussian model

$$P(z_t | x_t) = \mathcal{N}(z_t; x_t, \sigma_Q)$$





$x_0$

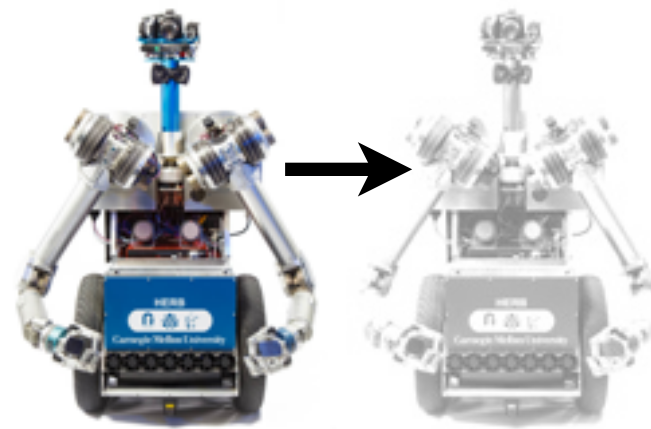
true position

$\hat{x}_0$

initial estimate

initial estimate uncertainty  $\sigma_0$

# Prior (initial) State



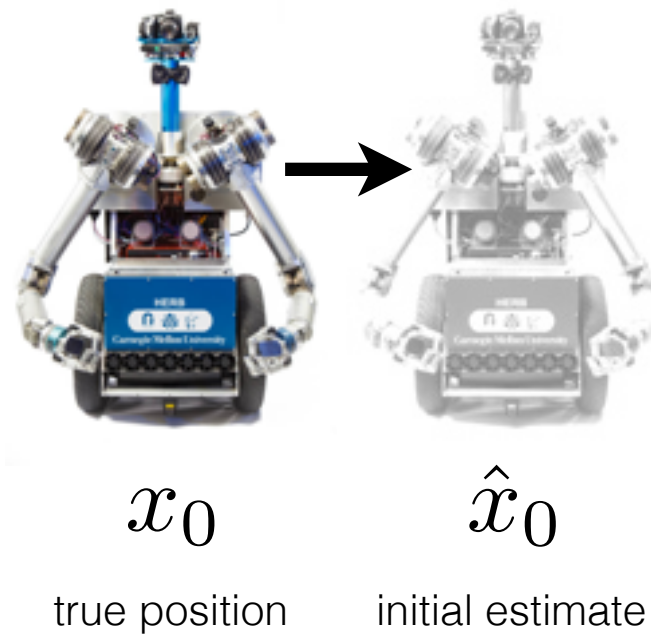
$x_0$

true position

$\hat{x}_0$

initial estimate

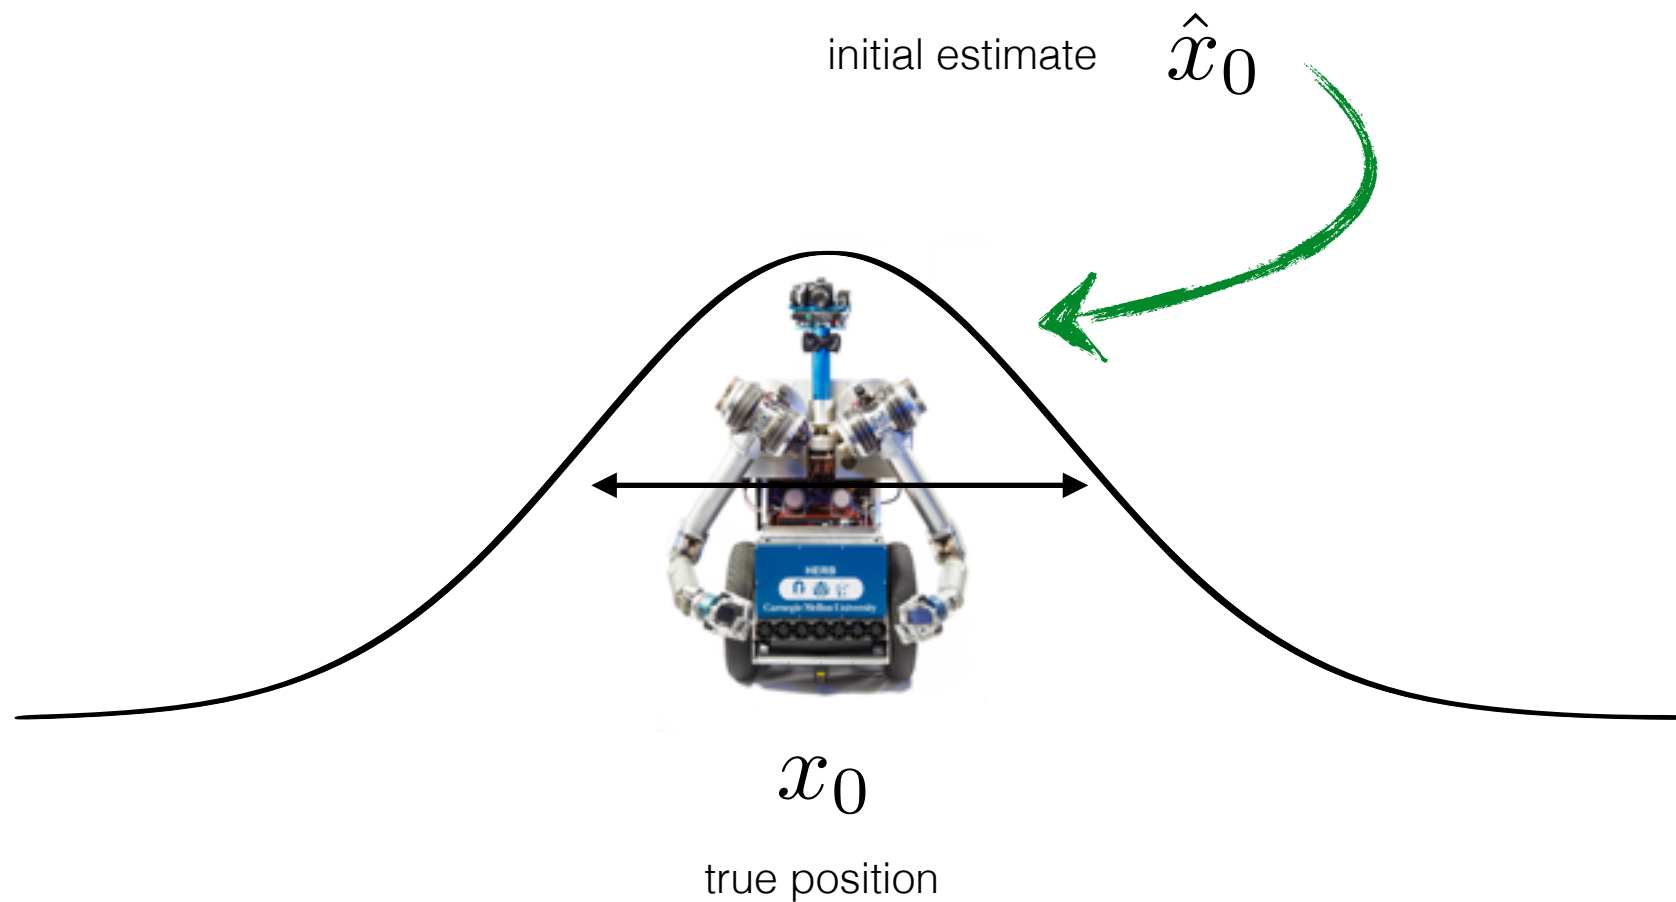
*How do you represent the prior state probability?*



*How do you represent the prior state probability?*

Also a linear Gaussian model!

$$P(\hat{x}_0) = \mathcal{N}(\hat{x}_0; x_0, \sigma_0)$$



*How do you represent the prior state probability?*

Also a linear Gaussian model

$$P(\hat{x}_0) = \mathcal{N}(\hat{x}_0; x_0, \sigma_0)$$

# Inference

*So how do you do temporal filtering with the KL?*

**Recall:** the first step of filtering was the ‘prediction step’

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \int_{\mathbf{x}_t} \underbrace{P(\mathbf{X}_{t+1} | \mathbf{x}_t)}_{\text{motion model}} \underbrace{P(\mathbf{x}_t | \mathbf{e}_{1:t})}_{\text{belief}} d\mathbf{x}_t$$

prediction step

compute this!  
It's just another Gaussian



need to compute the ‘prediction’ mean and variance...

# Prediction

(Using the motion model)

*How would you predict  $\hat{x}_1$  given  $\hat{x}_0$  ?*

using this 'cap' notation to  
denote 'estimate'

$$\hat{x}_1 = \hat{x}_0 + s \quad (\text{This is the mean})$$

$$\sigma_1^2 = \sigma_0^2 + \sigma_r^2 \quad (\text{This is the variance})$$

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \int_{\mathbf{x}_t} \underbrace{P(\mathbf{X}_{t+1} | \mathbf{x}_t)}_{\text{motion model}} \underbrace{P(\mathbf{x}_t | \mathbf{e}_{1:t})}_{\text{belief}} d\mathbf{x}_t$$

prediction step

the second step after prediction is ...



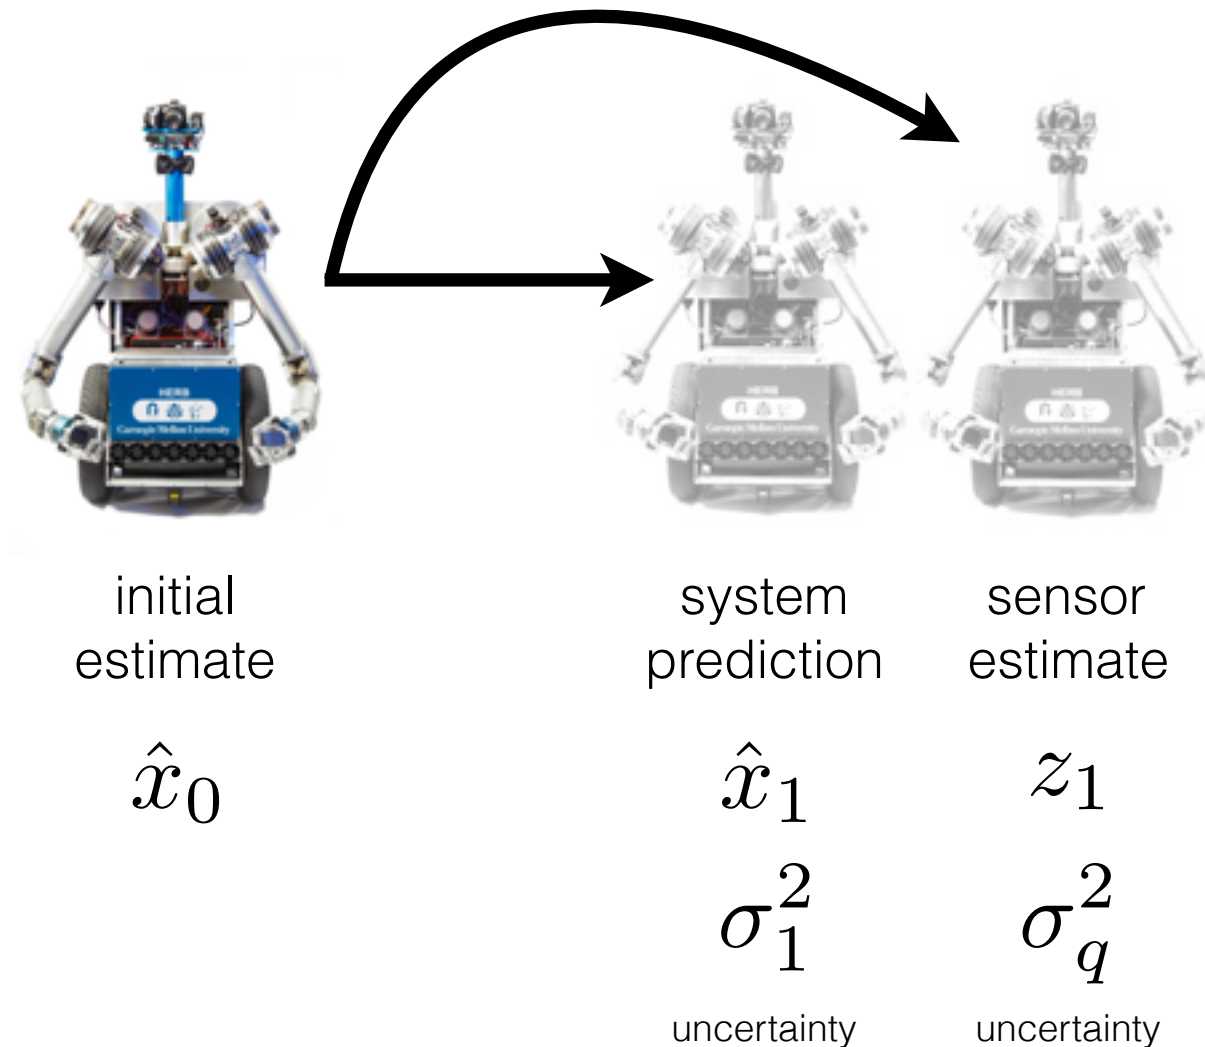
... update step!

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto \underbrace{P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1})}_{\text{observation}} \int_{\mathbf{x}_t} \underbrace{P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t})}_{\text{prediction}} d\mathbf{x}_t$$

compute this  
(using results of the prediction step)



In the **update step**, the **sensor measurement** corrects the system **prediction**



*Which estimate is correct? Is there a way to know?*

*Is there a way to merge this information?*

***Intuitively***, the smaller variance mean less uncertainty.



system  
prediction  $\sigma_1^2$



sensor  
estimate  $\sigma_q^2$

So we want a weighted state estimate correction



something  
like this...

$$\hat{x}_1^+ = \frac{\sigma_q^2}{\sigma_1^2 + \sigma_q^2} \hat{x}_1 + \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2} z_1$$

***This happens naturally in the Bayesian filtering (with Gaussians) framework!***

Recall the filtering equation:

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto \overbrace{P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1})}^{\text{observation}} \overbrace{\int_{\mathbf{x}_t} P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t}) d\mathbf{x}_t}^{\text{one step motion prediction}}$$

 Gaussian                       Gaussian

*What is the product of two Gaussians?*

Recall ...

When we multiply the prediction (Gaussian) with the observation model (Gaussian) we get ...

... a product of two Gaussians

$$\mu = \frac{\mu_1 \sigma_2^2 + \mu_2 \sigma_1^2}{\sigma_2^2 + \sigma_1^2}$$

$$\sigma = \frac{\sigma_1^2 \sigma_2^2}{\sigma_1^2 + \sigma_2^2}$$

applied to the filtering equation...

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \int_{\mathbf{x}_t} P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t}) d\mathbf{x}_t$$

mean:  $z_1$   
variance:  $\sigma_q$

mean:  $\hat{x}_1$   
variance:  $\sigma_1$

new mean:

new variance:

$$\hat{x}_1^+ = \frac{\hat{x}_1 \sigma_q^2 + z_1 \sigma_1^2}{\sigma_q^2 + \sigma_1^2}$$

$$\hat{\sigma}_1^{2+} = \frac{\sigma_q^2 \sigma_1^2}{\sigma_q^2 + \sigma_1^2}$$

'plus' sign means post  
'update' estimate



system  
prediction  $\sigma_1^2$



sensor  
estimate  $\sigma_q^2$

With a little algebra...

$$\hat{x}_1^+ = \frac{\hat{x}_1 \sigma_q^2 + z_1 \sigma_1^2}{\sigma_q^2 + \sigma_1^2} = \hat{x}_1 \frac{\sigma_q^2}{\sigma_q^2 + \sigma_1^2} + z_1 \frac{\sigma_1^2}{\sigma_q^2 + \sigma_1^2}$$

We get a weighted state estimate correction!

# Kalman gain notation

With a little algebra...

$$\hat{x}_1^+ = \hat{x}_1 + \frac{\sigma_1^2}{\sigma_q^2 + \sigma_1^2} (z_1 - \hat{x}_1) = \hat{x}_1 + \underset{\substack{\uparrow \\ \text{'Kalman gain'}}}{K} (\underset{\substack{\uparrow \\ \text{'Innovation'}}}{z_1 - \hat{x}_1})$$

With a little algebra...

$$\sigma_1^+ = \frac{\sigma_1^2 \sigma_q^2}{\sigma_1^2 + \sigma_q^2} = \left( 1 - \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2} \right) \sigma_1^2 = (1 - \mathbf{K}) \sigma_1^2$$



# Summary (1D Kalman Filtering)

To solve this...

$$P(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) \propto P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \int_{\mathbf{x}_t} P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t}) d\mathbf{x}_t$$

Compute this...

$$\hat{x}_1^+ = \hat{x}_1 + \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2} (z_1 - \hat{x}_1) \quad \sigma_1^{2+} = \sigma_1^2 - \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2} \sigma_1^2$$

$$K = \frac{\sigma_1^2}{\sigma_1^2 + \sigma_q^2}$$

‘Kalman gain’

$$\hat{x}_1^+ = \hat{x}_1 + K(z_1 - \hat{x}_1)$$

mean of the new Gaussian

$$\sigma_1^{2+} = \sigma_1^2 - K \sigma_1^2$$

variance of the new Gaussian

# Simple 1D Implementation

$$[x \ p] = KF(x, v, z)$$

$$x = x + s;$$

$$v = v + r;$$

$$K = v / (v + q);$$

$$x = x + K * (z - x);$$

$$p = v - K * v;$$

Just 5 lines of code!

or just 2 lines

$$\begin{aligned} [x \ P] &= KF(x, v, z) \\ x &= (x+s) + (v+r) / ((v+r)+q) * (z - (x+s)) ; \\ p &= (v+r) - (v+r) / ((v+r)+q) * v ; \end{aligned}$$

# Bare computations (algorithm) of Bayesian filtering:

KalmanFilter( $\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$ )

prediction mean  $\bar{\mu}_t = A_t \overset{\text{motion}}{\mu_{t-1}} + B \overset{\text{control}}{u_t}$  'old' mean Prediction

prediction covariance  $\bar{\Sigma}_t = A_t \overset{\text{'old' covariance}}{\Sigma_{t-1}} A_t^\top + R$  Gaussian noise

$K_t = \bar{\Sigma}_t C_t^\top (C_t \bar{\Sigma}_t C_t^\top + Q_t)^{-1}$  Gain

update mean  $\mu_t = \bar{\mu}_t + K_t (z_t - \overset{\text{observation model}}{C_t \bar{\mu}_t})$

update covariance  $\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$  Update

# Simple Multi-dimensional Implementation (also 5 lines of code!)

$$[x \ P] = KF(x, P, z)$$

$$x = A * x;$$

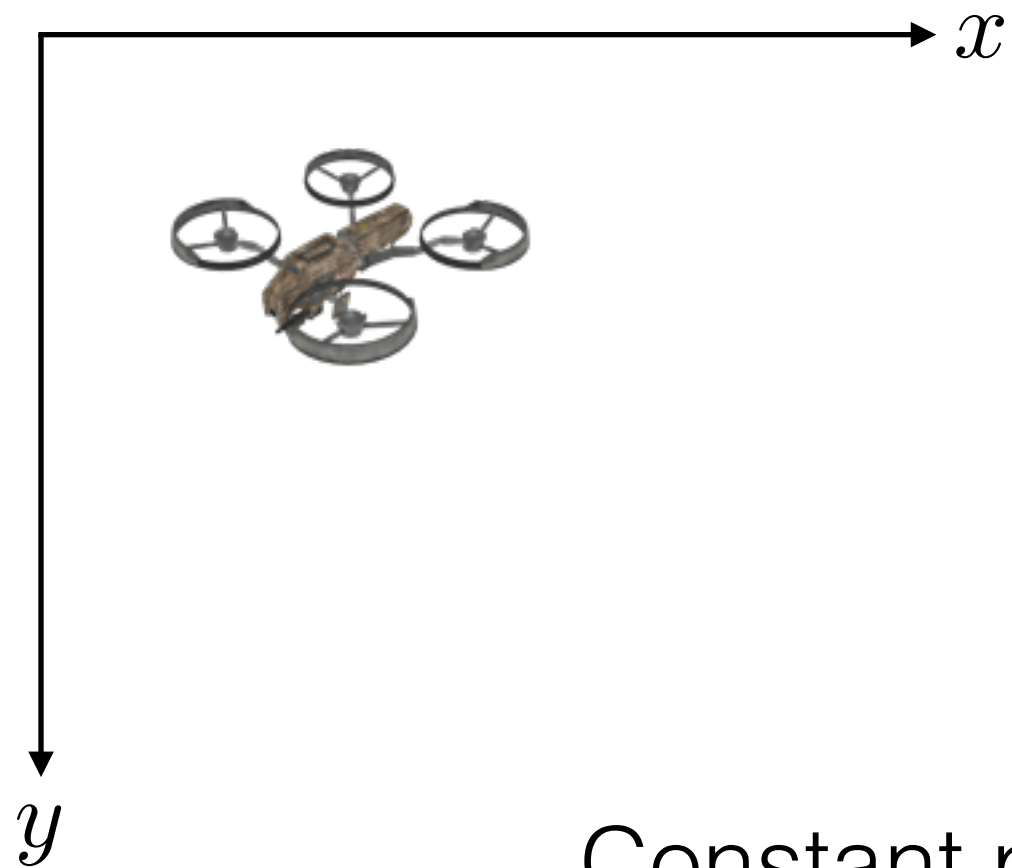
$$P = A * P * A' + R;$$

$$K = P * C' / (C * P * C' + Q);$$

$$x = x + K * (z - C * x);$$

$$P = (\text{eye}(\text{size}(K, 1)) - K * C) * P;$$

# 2D Example



state

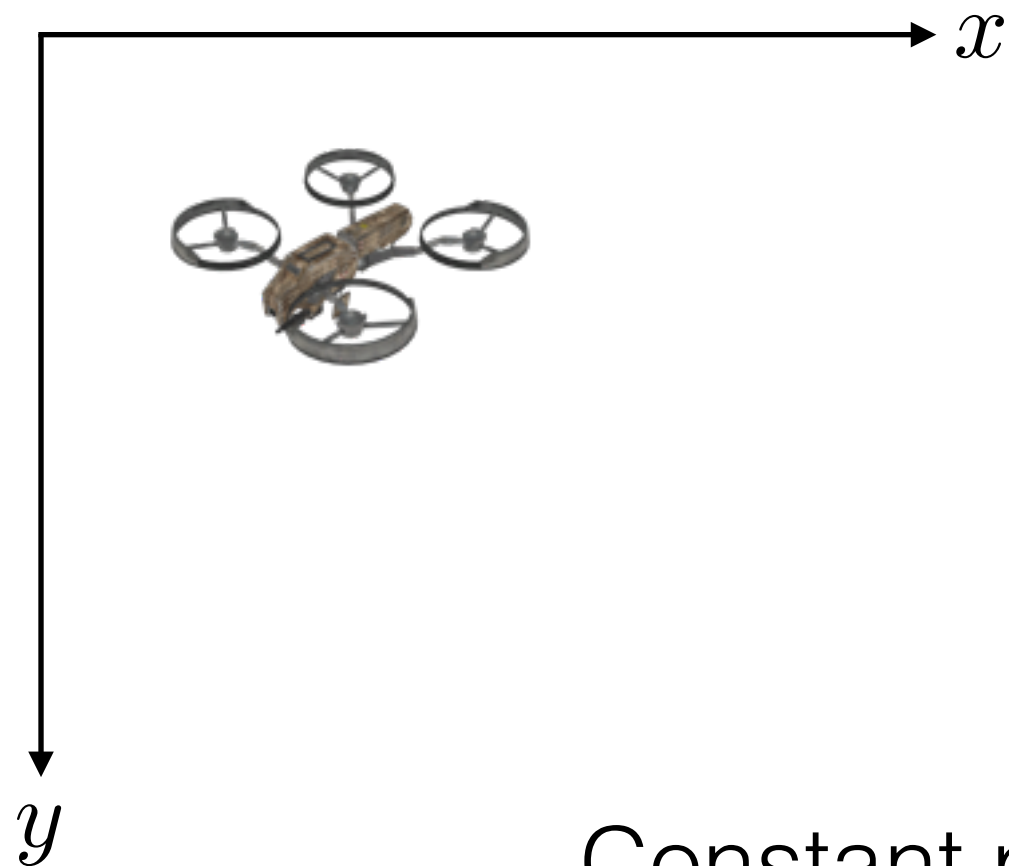
$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Constant position Motion Model

$$\mathbf{x}_t = A\mathbf{x}_{t-1} + B\mathbf{u}_t + \epsilon_t$$



state

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Constant position Motion Model

$$\mathbf{x}_t = A\mathbf{x}_{t-1} + B\mathbf{u}_t + \epsilon_t$$

system noise

$$\epsilon_t \sim \mathcal{N}(\mathbf{0}, R)$$

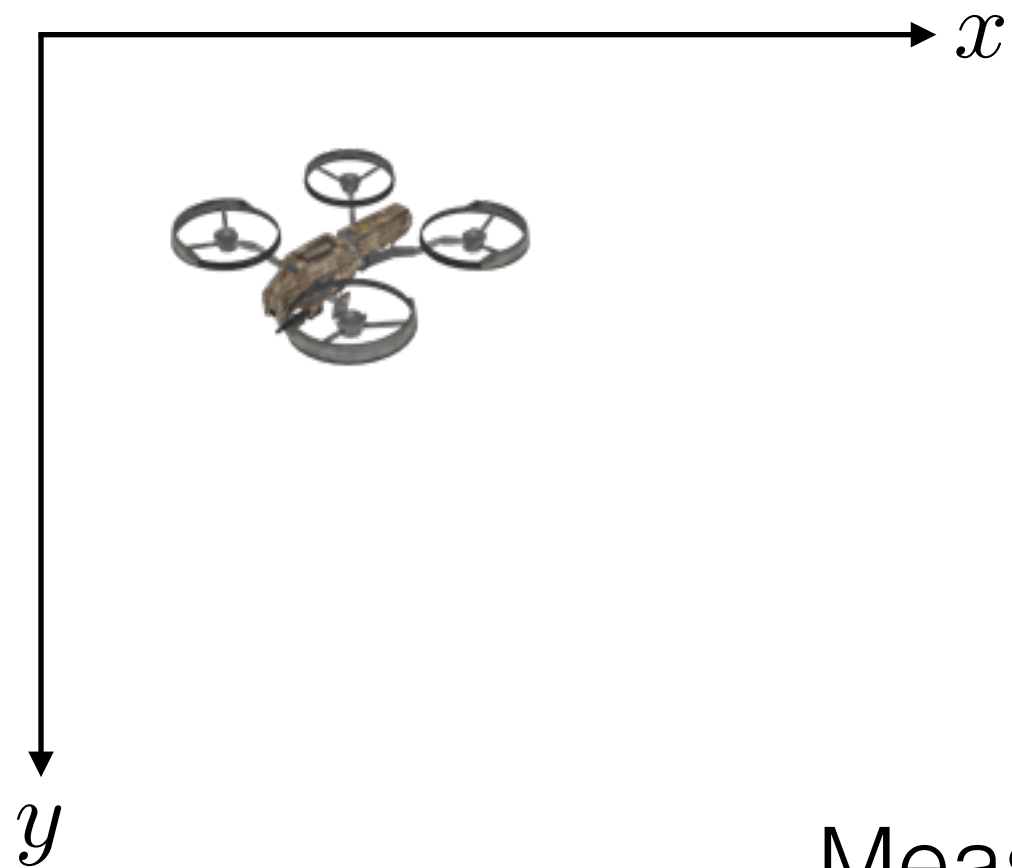
Constant position

$$A = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$B\mathbf{u} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$R = \begin{bmatrix} \sigma_r^2 & 0 \\ 0 & \sigma_r^2 \end{bmatrix}$$





state

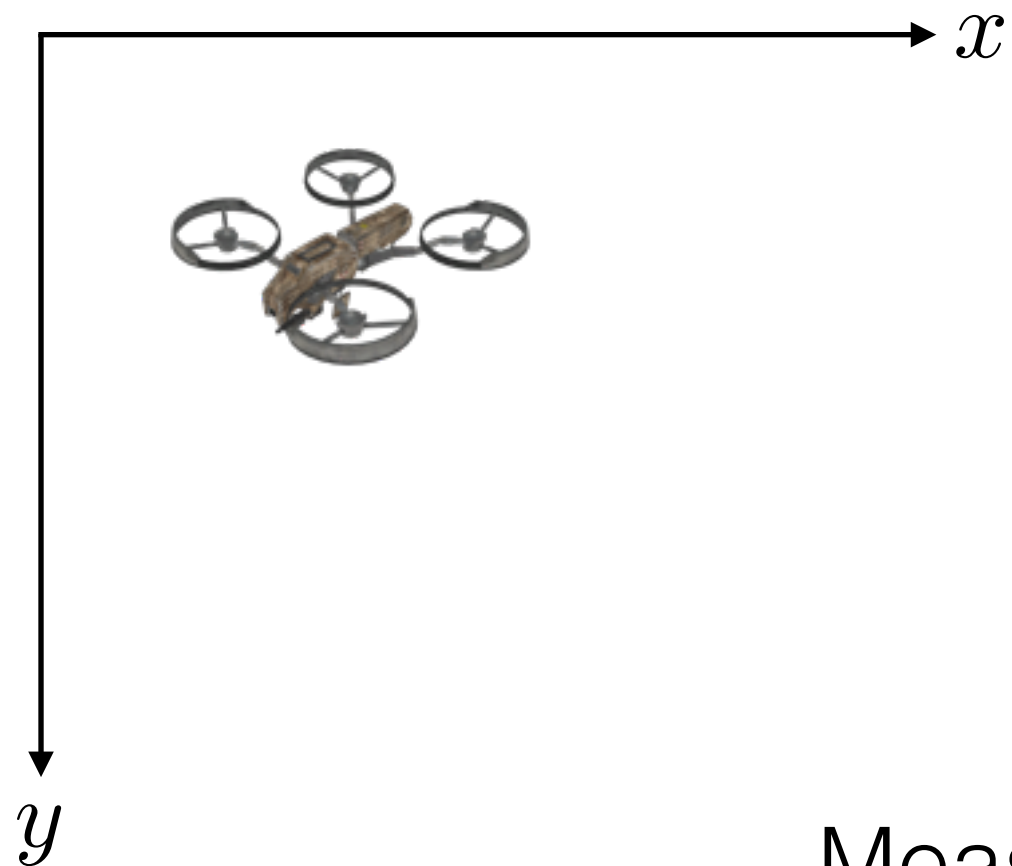
$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Measurement Model

$$\mathbf{z}_t = \mathbf{C}_t \mathbf{x}_t + \delta_t$$



state

$$\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$$

measurement

$$\mathbf{z} = \begin{bmatrix} x \\ y \end{bmatrix}$$

Measurement Model

$$\mathbf{z}_t = \mathbf{C}_t \mathbf{x}_t + \delta_t$$

zero-mean measurement noise

$$\mathbf{C} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$\delta_t \sim \mathcal{N}(\mathbf{0}, \mathbf{Q})$$

$$\mathbf{Q} = \begin{bmatrix} \sigma_q^2 & 0 \\ 0 & \sigma_q^2 \end{bmatrix}$$

# Algorithm for the 2D object tracking example



$$A = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

motion model

$$C = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

observation model

## General Case

$$\bar{\mu}_t = A_t \mu_{t-1} + B u_t$$

$$\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^\top + R$$

$$K_t = \bar{\Sigma}_t C_t^\top (C_t \bar{\Sigma}_t C_t^\top + Q_t)^{-1}$$

$$\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$$

$$\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$$

## Constant position Model

$$\bar{\mathbf{x}}_t = \mathbf{x}_{t-1}$$

$$\bar{\Sigma}_t = \Sigma_{t-1} + R$$

$$K_t = \bar{\Sigma}_t (\bar{\Sigma}_t + Q)^{-1}$$

$$\mathbf{x}_t = \bar{\mathbf{x}}_t + K_t (z_t - \bar{\mathbf{x}}_t)$$

$$\Sigma_t = (I - K_t) \bar{\Sigma}_t$$

Just 4 lines of code

```
[x P] = KF_constPos(x, P, z)
```

```
P = P + R;
```

```
K = P / (P + Q);
```

```
x = x + K * (z - x);
```

```
P = (eye(size(K,1)) - K) * P;
```

*Where did the 5th line go?*

## General Case

$$\bar{\mu}_t = A_t \mu_{t-1} + B u_t$$

$$\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^\top + R$$

$$K_t = \bar{\Sigma}_t C_t^\top (C_t \bar{\Sigma}_t C_t^\top + Q_t)^{-1}$$

$$\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$$

$$\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$$

## Constant position Model

$$\bar{\mathbf{x}}_t = \mathbf{x}_{t-1}$$

$$\bar{\Sigma}_t = \Sigma_{t-1} + R$$

$$K_t = \bar{\Sigma}_t (\bar{\Sigma}_t + Q)^{-1}$$

$$\mathbf{x}_t = \bar{\mathbf{x}}_t + K_t (z_t - \bar{\mathbf{x}}_t)$$

$$\Sigma_t = (I - K_t) \bar{\Sigma}_t$$